


We don't build demos.

We build production systems that run at scale, with governance and measurability built in.

END-STATE ARCHITECTURE

A single, client-owned orchestration layer between enterprise systems and the model layer. Every AI action passes a confidence gate; every human decision is recorded.



WHERE COMMERCE BREAKS DOWN

Beauty and luxury retail run on product data that lives in five systems and reconciles in none. Every problem below shows up as a confidence gap somewhere downstream.

- 01 Product data fragmented across regions & channels**
HQ, regional teams, marketplaces, and retail partners each maintain their own version. Reconciliation happens manually, country by country.
- 02 Multilingual content can't keep pace with launches**
Translation and localization lag behind launches. Brand voice drifts between markets. Translation memory and glossary live in disconnected tools.
- 03 AI rewrites ship without governance**
Generated copy goes live without similarity gates, glossary checks, or tone review. Retraction in a luxury context costs more than the original work.
- 04 Channel adaptations rebuilt by hand**
Each marketplace, retailer feed, and DTC site needs its own copy and asset variants. Nothing reuses prior work or learns from cycles.

THE PROOF

What the architecture delivers, measured. Every number below is from a system running in production today.

42→20

API surface reduced post-UAT

Engineering simplicity as a deliverable.

10

Governed workflow states

Full state machine, zero manual tracking.

5

Enterprise systems integrated

One governed pipeline, end-to-end.

8–12 wks

Discovery to production

Gated delivery, signed off at every phase.

100%

Audit coverage

Every AI call, every human decision, every action.

0

Ungoverned AI outputs

Confidence gate before any downstream action.

HOW WE ENGAGE

A gated delivery model. Every phase ends in a signed-off demo before the next one starts.

01 WK 1–2
Discovery
Joint working sessions with your team. We end with a written architecture, a working evaluation harness, and a go / no-go.

02 WK 2–4
Vertical slice
One end-to-end path through the system, from capture to governed output, wired to real enterprise systems, not mocks.

03 WK 4–8
Build & harden
Scope expands across agents, integrations, and governance. Evals run on every change; nothing ships without a passing gate.

04 WK 8+
Operate
You own the system. We hand over runbooks, monitoring, and an evaluation pipeline; we stay on as engineering reserve.