THE FABLE 5 LESSON

If your LLM disappeared tomorrow, would your clients notice before you did?

NEWTUPLE

WHAT FABLE 5 EXPOSED

You don't own your AI product. You're renting its most important part.

Coding agents, internal tools, client deliverables, decision pipelines, all sitting on a model someone else controls.

With Fable 5, that model went away, and the work built on top of it went with it.

THE DRIFT NOBODY SEES

You didn't change a thing. The model did.

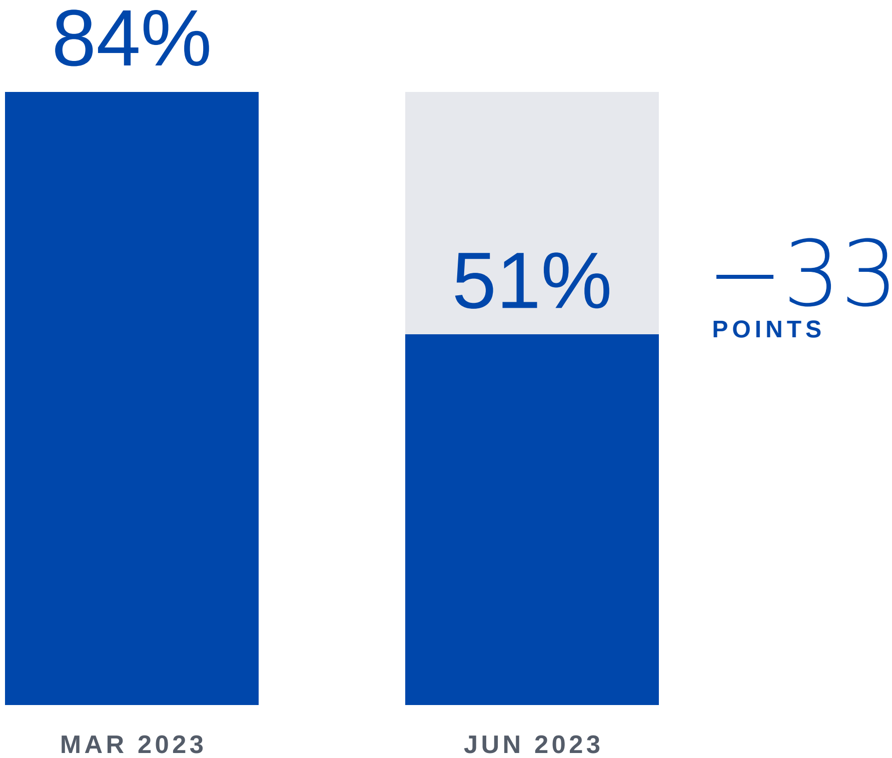

GPT-4 on the same prime-number test, three months apart.

SOURCE: CHEN, ZAHARIA & ZOU · STANFORD/UC BERKELEY · ARXIV 2307.09009

AND IT IS SILENT

Nobody sends a changelog when your model gets worse.

The same researchers found that when and how these services update is opaque, no notice, no diff, no warning.

Your most important dependency can change behavior between two ordinary Tuesdays.

THE REFRAME

This is not an AI problem. It is an eval problem.

The model leaving only hurt because nothing measured its quality on your work, nothing watched for the drop, and nothing had a tested replacement ready to go.

HOW MODELS GET CHOSEN

You're picking a production dependency off a launch-day screenshot.

Most model decisions still come down to:

- Benchmark screenshots
- Launch-day hype
- A thread that went viral
- The vendor's own numbers
- “It feels better when I use it”

WHAT ACTUALLY MATTERS

Production doesn't care how it felt in the demo.

It asks:

- Does it work on our tasks?
- Does it fail safely?
- Can we catch regressions early?
- Can another model drop straight in?
- Are we tracking quality, latency, and cost?
- Do we know when behavior shifts?

THE FIX

Stop building dependencies on someone else's model.

GAUGETUPLE · NEWTUPLE'S EVALUATION ACCELERATOR

Benchmark and compare LLMs on the workflows that actually run your business, catch regressions before clients do, and prove a model works on your own cases before you ship.

BENCHMARK > EVALUATE > PROVE

Benchmark your model against its closest alternatives.

Line up the model you run today against the nearest alternatives and see which one wins on your own workflows — quality, latency, and cost, side by side.

- Your current model vs. the closest alternatives
- Quality, latency, and cost in one ranked view
- Scored on your tasks, not a vendor's benchmark

dev.gaugetuple.com/benchmark

Model Benchmark

Support-resolution workflow · ranked on 1,240 of your tasks

MODEL	ACCURACY	LATENCY	COST/1K
Claude Sonnet 4.5 BEST	**94.2%**	1.1s	$3.0
GPT-4o	**91.8%**	0.9s	$5.0
Gemini 2.5 Pro	**89.5%**	1.4s	$4.0
GPT-4 (current) CURRENT	**87.1%**	1.2s	$10
Llama 3.1 70B	**82.3%**	0.7s	$0.6

INSIDE GAUGETUPLE · EVALUATE

Run constant evals on the questions your users actually ask.

Score every model answer against your golden answers with your own evaluation prompts — and tighten prompts on evidence instead of gut feel.

- Real user questions vs. your golden answers
- Pass / fail scoring with your own eval prompts
- Catch the 4% that fails before it ships

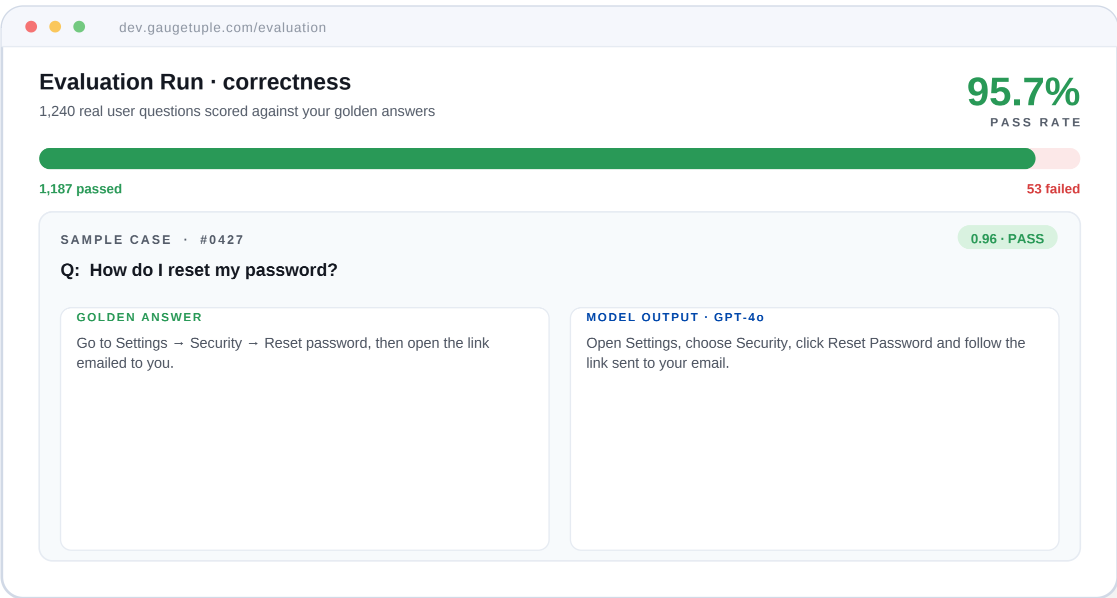

Catch the silent drift before your clients do.

Constant evals keep running as the model underneath you shifts. The moment quality slips, GaugeTuple flags it — with a tested replacement ready to drop in.

- Continuous monitoring across model versions
- Alerts the moment behavior regresses
- A tested, drop-in replacement on standby

THE REAL WARNING

What happens to your AI when the model underneath it changes?

Fable 5 was a preview, not an exception. If your product leans on one model, one API, one vendor's benchmark, you are one shutdown away from explaining failure to a client.

BENCHMARK IT. EVALUATE IT. PROVE IT.

NEWTUPLE